

CS307

Database Principles

Stéphane Faroult
faroult@sustc.edu.cn

朱悦铭 Zhu Yueming zhuym@sustc.edu.cn

What we have seen last time:

Indexes

B-Trees – **VERY** important for performance


PK and unique constraints create indexes

Has to be maintained / storage


Cannot always be used

EXPLAIN

Indexing expressions



Don't fall into the trap that says "if I have a query that is slow, I just have to add an index and it will be fast". The index HAS TO BE USABLE, and besides remember that it depends on volumes returned.



In fact, proper index usage depends a lot on how you write your queries and this is what we are going to see now.



We have seen that joins and subqueries could be exchanged. But all writings haven't the same indexing implications.

Let's take this join. Basically we can run it in two different ways, better explained with subqueries.

```

select count(c.movieid)
from credits c
  inner join people p
    on p.peopleid = c.peopleid
where c.credited_as = 'A'
   and p.surname = 'Hanks'
   and p.first_name = 'Tom'

```

We can look for all peopleid values for actors, and check whether it's Tom Hanks'

```

select count(c.movieid)
from credits c
where c.credited_as = 'A'
   and exists (select null
               from people p
               where p.peopleid = c.peopleid
                  and p.surname = 'Hanks'
                  and p.first_name = 'Tom')

```

Or we can look for Tom Hanks' peopleid value, and look for films in which it's found with 'A'

```

select count(c.movieid)
from credits c
where c.credited_as = 'A'
   and c.peopleid in (select peopleid
                     from people p
                     where p.surname = 'Hanks'
                        and p.first_name = 'Tom')

```

It means completely different things. In the first case I scan CREDITS and use the PK in PEOPLE each time

```

select count(c.movieid)
from credits c
where c.credited_as = 'A'
   and exists (select null
               from people p
               where p.peopleid = c.peopleid
                  and p.surname = 'Hanks'
                  and p.first_name = 'Tom')

```

In the second case I use the PEOPLE index on (SURNAME, FIRST_NAME)

```

select count(c.movieid)
from credits c
where c.credited_as = 'A'
   and c.peopleid in (select peopleid
                     from people p
                     where p.surname = 'Hanks'
                        and p.first_name = 'Tom')

```

I may not be able to use the PK on CREDITS

(movieid, **peopleid**, credited_as)

Constraints with an eye to Performance

Which brings us to another topic: when I define constraints, can I do it in a clever way?

Primary key

INDEX

Unique constraint

I have told you that PK and UNIQUE constraints were both creating an index, to quickly check whether an entry is already known.

UNIQUE

(first_name, surname)

(surname, first_name)

Which one?

From a logical point of view, when you say that a combination of columns is unique, the order doesn't matter. If it doesn't matter, perhaps we can use it at our benefit?

I can expect different types of queries that will refer to the columns in the unique constraint.

```
select *
from people
where surname = 'Spielberg'
and first_name = 'Steven'
```

```
select *
from people
where surname = 'Spielberg'
```

```
select *
from people
where first_name = 'Steven'
```

In all likelihood the most common ones will be the first two ones. If I want the index to be usable in both cases SURNAME should come first.

```
select *
from people
where surname = 'Spielberg'
and first_name = 'Steven'
```

```
select *
from people
where surname = 'Spielberg'
```

```
select *
from people
where first_name = 'Steven'
```

second

first

What about storing rows in some order?

Another idea to speed up queries would be to store rows in a given order. If order doesn't matter from a relational point of view, we can once again perhaps use it for our own benefit.

A regular table (unordered, rows are inserted where room is found) is called a

Heap-organized table

A table that is ordered is called a




Clustered index

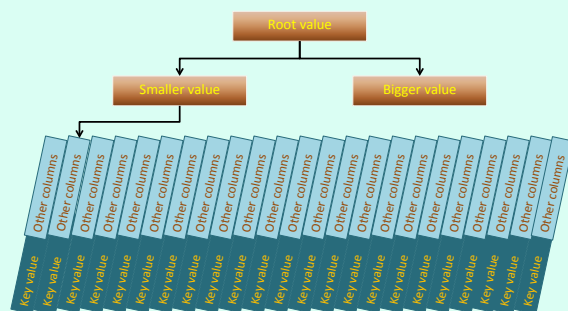
ORACLE Index-organized table

why "index"? Because indexes are strongly ordered structures, and an ordered table follows the same pattern.

In an index, you find key values and row addresses.




In an ordered table, row data replaces the row address.



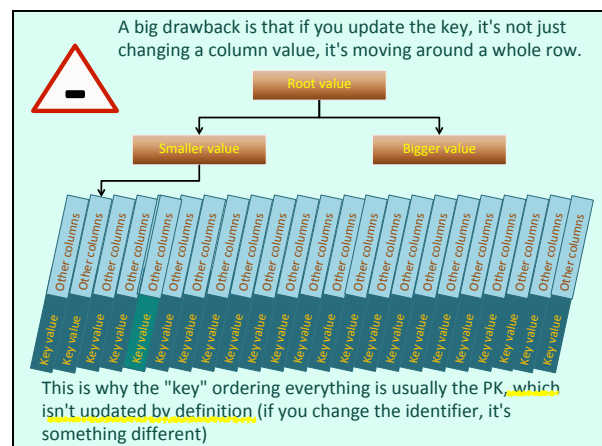
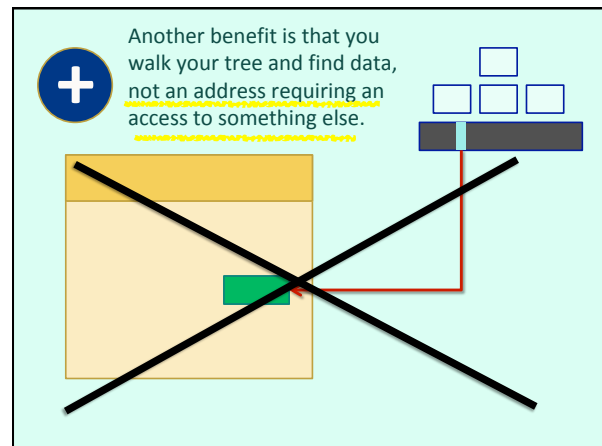
For internal storage reasons, it works better with few columns.

That way you have all rows ordered according to the key value.



Benefits Drawbacks

Like any clever idea, it has its benefits and drawbacks. Interestingly, this structure is much used with SQL Server (almost standard) and with MySQL/InnoDB, and rarely with Oracle.





insert

Inserting a row also means inserting it at the right place, which may involve shifting around a lot of bytes to make room. Not an issue if rows are inserted in the same order as the key (for instance if the key is an auto-numbered column or the timestamp at insertion). Big issue if the order of insertion is more or less random in reference to the key.



Right place

Natural candidates for this type of organization is tables in the "tall and thin" (many rows, few columns) category, not in the "short and plump" one, and for which the key order makes some real-life sense for range scans.

Many rows

Few columns

Tables that implement a many-to-many relationship may fall into this category.

Natural order

movies

movieid

In the film database, the order of rows makes no particular sense for MOVIES or PEOPLE. There is no reason why we would be interested by ranges of films or person identifiers.

people

peopleid

credits

~~where movieid between 350 and 427~~

CREDITS is a different case. Its key is the concatenation of MOVIEID, PEOPLEID and CREDITED_AS. If we look for a MOVIEID value, it will be a range scan (of all keys starting with this value). Besides, it's useless to have an index that stores exactly the same

information as the table, plus locators.



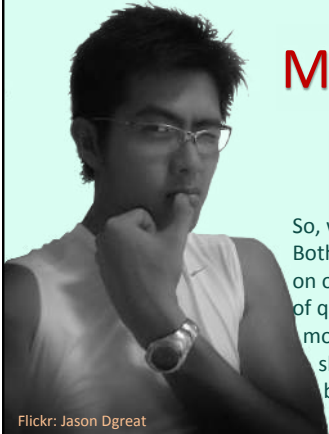
And here, it makes some real-life sense of finding people involved in a film clustered together.

movieid	peopleid	credited_as
123	178	D
123	312	A
123	86	A
123	105	A
123	237	A



However, it would also make some real-life sense of finding all the films of one person clustered together.

peopleid	movieid	credited_as
97	178	D
97	178	A
97	212	D
97	212	A
97	126	A

Movies ?

People ?

So, what should we put first? Both make sense. It depends on our audience, and the type of query that we expect to be most frequent. But in fact we shouldn't worry too much, because we can have our cake and eat it.

Flickr: Jason Dgreat

Index both!
(constraint + additional index)
Order by the most important for you

```
main() {
  my_func()

  my_func()

  my_func()
}
```

```
my_func() {
  
```

We have seen stored functions, but these functions are just returning numbers, strings or dates. Can we have relational functions, allowing to reuse relational operations as we would reuse code in a program?



**create view *viewname* (col1, ..., coln)
as
select ...**

In practice (theory is a bit more complicated) there isn't much to a view: it's basically a named query. If the query is correct, it should return a valid relation, so why not consider it as if it were a table? You can optionally rename columns after the view name (if you don't, the view uses column names from the query result)


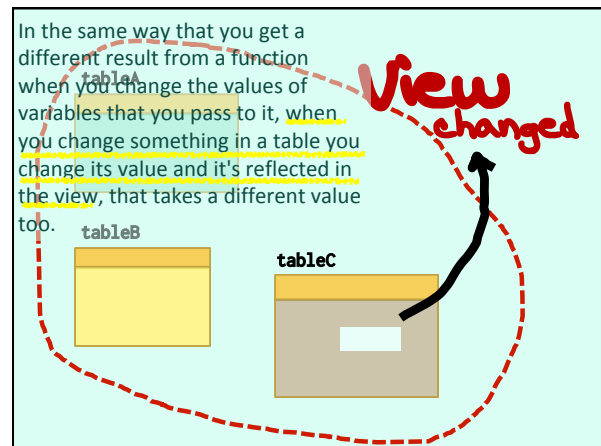

tableA

tableB

tableC

View

A view can be as a complicated query as you want, and will usually return something that isn't as normalized as your tables, but easier to understand.



```
create view vmovies
as select m.movieid,
       m.title,
       m.year_released,
       c.country_name
from movies m
inner join countries c
on c.country_code = m.country
```

```
select *
from vmovies
where country_name = 'Italy'
```

Once the view is created, I can query the view exactly as if it were a table; nothing says that it's a view, except the name that *I* have chosen. I like to give a special name to views to make it clear that it's a view (discussion about practical differences between views and tables comes soon) but I could have masked a change in table design to allow old programs to run by having a changed table T renamed T_V2 and creating a view T rendering the old version.

```
select *
from (select m.movieid,
           m.title,
           m.year_released,
           c.country_name
       from movies m
       inner join countries c
         on c.country_code = m.country)
vmovies
where country_name = 'Italy'
```

Result-wise, the previous query is strictly equivalent to this one.

Some optimizers are able to push the condition up into the view.

However, there is far more than this to views. As I have said earlier, views are just the relational equivalent of functions: the ability to store (and reuse) a relational expression, in other words something that returns a relation and not simply a value like what you usually do with a stored function.

If we step back to design issues, you remember that modelling a database is basically distributing data between normalized tables, and there are often ways of organizing data that are more suitable for a given application. In some respect, views provide a way of creating an "alternate model".

What is important is that views are permanent objects in the database - needless to say, their content will change with the data in the underlying tables, but the structure will remain constant and can be described in the same way as the structure of a table can be described: columns are typed. Beware that columns are the one in tables when the view was created. Columns added later to tables in the view won't be added even if the view was created with SELECT * (bad practice)

Permanent object
Permanent structure



In real life, views are much used for simplifying queries. Many business reports are based on the same set of hairy joins, with just variations on the columns that you aggregate or order by. Somehow, views allow to come back to that old fancy of the early days of SQL, having something that anybody can query with simple commands, without having to master the intricacies of the querying language.

Simplify queries


Make the sharp guys write the hard stuff, then use cheap code-monkeys to wrap some basic things around... That's more or less the idea (but it's rarely presented like this).

**Leverage
the skills
of
THE BEST
SQL CODERS**

```
create view vmovie_ere
as select m.title,
m.year, m.cla,
cast(c.cla as
when 'A' then
when 'D' then 'D'
else '?'
end role,
f.fname(p),
p.surname) name
from movies m
inner join c
on c.m = m
inner join p
```

```
select *
from vmovies
where country_name = 'Italy'
```

This is something that a cheap beginner completely ignorant of databases should be able to write after having been briefed for about three minutes.



Looks like a table
Tastes like a table

But

It's again an area where theory is bruised by practice.



```
create view vmovie_credits
as select m.title,
       m.year_released release,
       case c.credited_as
         when 'A' then 'Actor'
         when 'D' then 'Director'
         else '?'
       end duty,
       full_name(p.first_name, p.surname) name
from movies m
  inner join credits c
    on c.movieid = m.movieid
  inner join people p
    on p.peopleid = c.peopleid
```

Let's say that we have this view, which nicely displays film credits, including people names like 'Erich von Stroheim' as they should appear.

vmovie_credits

title	release	duty	name
Casablanca	1942	Director	Michael Curtiz
Casablanca	1942	Actor	Humphrey Bogart
Casablanca	1942	Actor	Ingrid Bergman
Casablanca	1942	Actor	Conrad Veidt
Casablanca	1942	Actor	Claude Rains
...

When we query the view, it looks really good and user-friendly. Well, it actually depend on HOW we query it, and on which column. Querying by title will be fine.

ORACLE'

And sometimes the optimizer can do very clever things.

```
select country, count(*) as num_films
from movies
group by country
having country = 'ar'
```

I have told you that Oracle (at least the last time I checked) would, with this type of query, process the aggregate, then discard everything that isn't Argentine. It should be a WHERE condition.

ORACLE'

If you create this view

```
create view movie_count
select country, count(*) as num_films
from movies
group by country
```

and run this query you might expect the same phenomenon to occur

```
select *
from movie_count
where country = 'ar'
```

ORACLE'

In fact, no. The optimizer will see the problem.

```
select country, count(*) as num_films
from movies
where country = 'ar'
group by country
```

It will "push" the condition up into the query and only count Argentine films, running in effect the query above.

```
create view vmovie_credits
as select m.title,
```

```
    m.year_released release,
    case c.credited_as
      when 'A' then 'Actor'
      when 'D' then 'Director'
      else '?'
    end duty,
    full_name(p.first_name, p.surname) name
from movies m
inner join credits c
  on c.movieid = m.movieid
inner join people p
  on p.peopleid = c.peopleid
```

There are times, though, when all the benevolence of the optimizer cannot do

anything for you. You may remember how awful function full_name() is

```
select *
from vmovie_credits
where name = 'Humphrey Bogart'
```

Dreadful expression

If you are writing something like this, what looks like a column (NAME) is in fact the result of a function. There is no way the index on (SURNAME, FIRST_NAME) can be used. We'll have to scan the full table, compute the function, and compare its result to the constant. Unless you do some tricky stuff to index in a way or the other the result of the function (not always possible).

VIEWS

Hide complexity

The problem with views is that as long as you haven't seen how they have been defined, you have no idea how complex they may be. They may be fairly innocuous, or they may be queries of death (they often are)

```
select distinct title
from vmovie_credits
```

Difficulties usually increase sharply as a young developer gets with time more confident, not to say bold, with SQL. Being so accustomed to working with this convenient "table", VMOVIES_CREDITS (it may not bear a name that makes it obvious it's a view), the developer may think of this as a way to return all the different titles in the database. Technically speaking, it will return the desired result, and it may even do it reasonably fast.

```
select distinct title
from
  (select m.title,
    m.year_released release,
    case c.credited_as
      when 'A' then 'Actor'
      when 'D' then 'Director'
      else '?'
    end duty,
    full_name(p.first_name, p.last_name) name
  from movies m
   inner join credits c
    on c.movieid = m.movieid
   inner join people p
    on p.peopleid = c.peopleid) vmovie_credits
```

Lot of useless work for what we want

Do we really want the join?

Scalability

And here we are coming to one of the great issues with databases and information systems generally speaking, namely the ability to deliver response times that remain acceptable when the number of users, data volumes, or both, sharply increase.

The computer system of any retailer must survive Black Friday in the US or 11/11 in China.

Computing power is always

LIMITED

And the problem is that it doesn't matter how big and powerful your computers are, computing power will always be a limited resource.

Slower query to retrieve the same data



Fewer simultaneous users served

You will only be able to serve *that* many users simultaneously.

You don't want to see everything crawl during peak time.

Query table





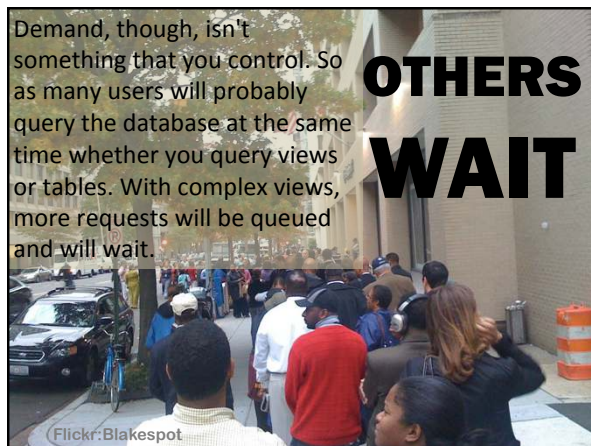
Query view



It's mechanical.

If querying the view takes 4 times as long as querying the table, you'll only be able to serve 25% of users..





Nevertheless, there are three areas where views are very useful. I have mentioned reporting, user interfaces are a bit in the same spirit (more later) the third area is security.

Reports
User Interface
SECURITY

TOP SECRET

How is
security
managed?

Before we see how views can help, we need to review how security is managed in a database.

To access a database, you must be authenticated, which often means entering a username and a password.

Username :

Password:


Connect

There are other means of authentication, and for some products database authentication is tied to operating system authentication, but in any case the database knows who you are.

So you end up being connected to a database account, and this account as a set of rights.

Database Account **RIGHTS**

In fact those rights are called "privileges", so that you understand that getting them is a favor.



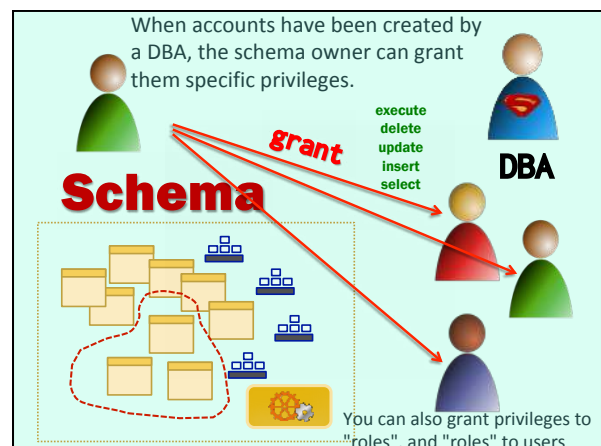
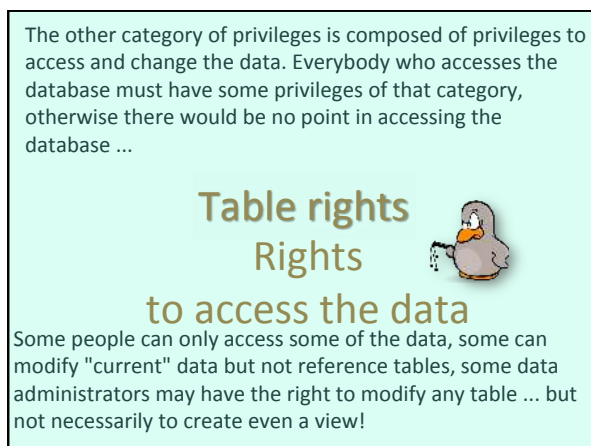
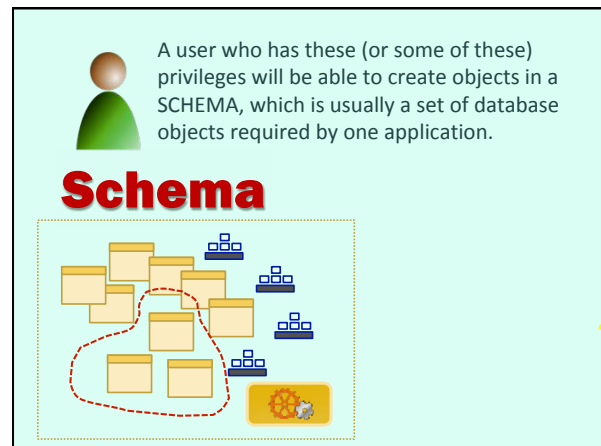
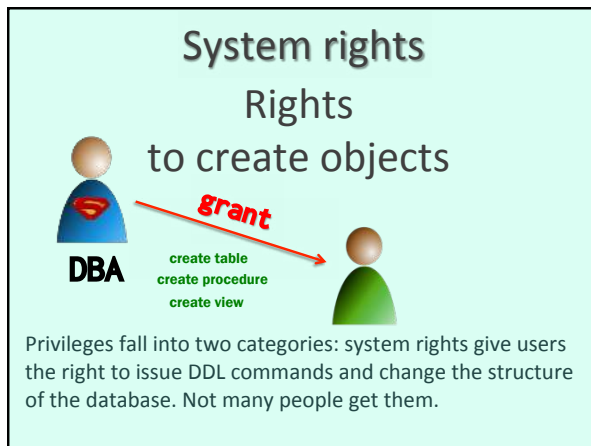
A privilege is given to a user account using this command:

grant *<right>* **to** *<account>*

and can be taken back using this one:

revoke *<right>* **from** *<account>*

GRANT and REVOKE are the two pillars of what is sometimes called DCL, Data Control Language.



GRANT commands to give privileges on a table look like this. You can give one or several privileges at once. Sometimes you can give privileges over all the tables in a schema, existing tables and tables still to be created. The UPDATE privilege can also be restricted to some columns only. Some products may require special additional rights (with PostgreSQL "usage" on a schema)

```
grant select, insert on tablename
to accountname
```

And for users who have been naughty:

```
revoke privilege on tablename
from accountname
```

So ...

How can views help with security?

The trick is to use a view that only shows what people are supposed to see, and grant SELECT on the view and not on the table..

people grant select on view

[illegible]

You can hide sensitive columns

Flickr: Nate Steiner

```
create view my_stuff
as
select * from stuff
where username = user
```

Syntax for identifying the current user varies.

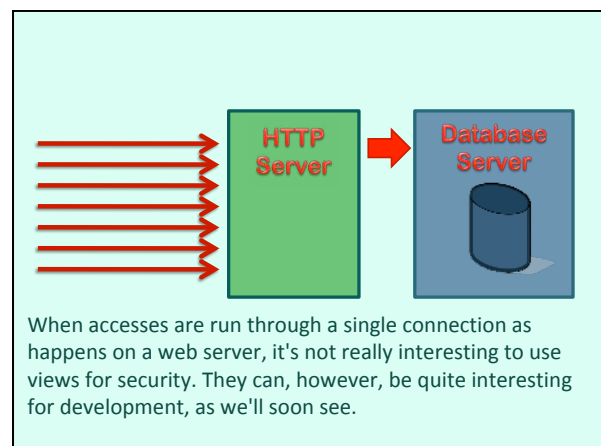
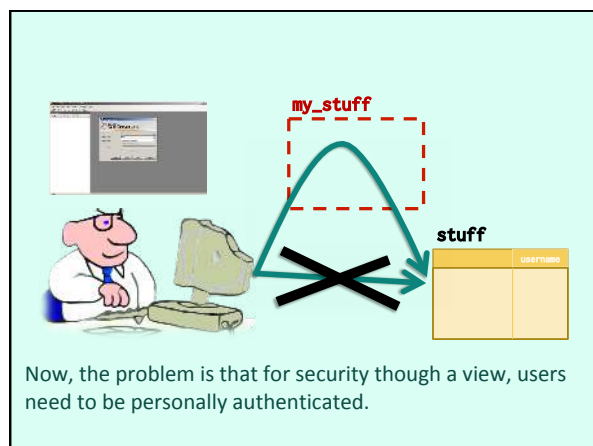
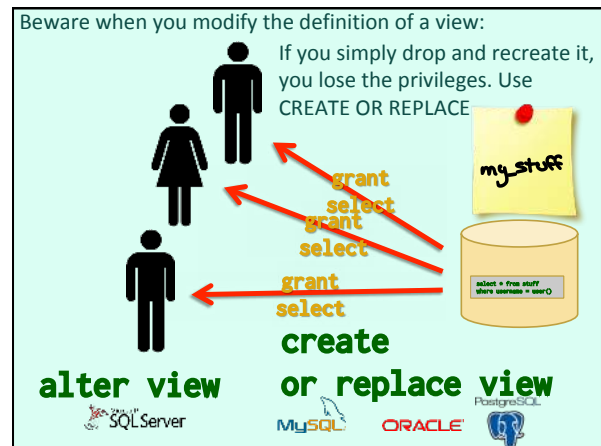




stuff

	username

You can even hide rows by only returning rows "owned" by the user currently connected.

my_stuff





**What about
CHANGING DATA
through views?**

If views are in theory like tables, why not using them for controlling not only what you SEE, but what you CHANGE?

Lots of things can go wrong

It all depends on the view ... The problem is that most view are designed to provide a more user-friendly view of data: joins transforming codes into more legible values, functions making data prettier (date formatting, for instance). And by doing so you often lose information.

For instance if your view concatenates first_name and surname, splitting a single string in two parts is tough if you want to insert through the view.

```
case
  when p.first_name is null then p.surname
  else p.first_name || ' ' || p.surname
end name
```

Tommy Lee Jones

Benicio Del Toro

Everybody isn't called 'Gary Cooper'.

And for updates ... Let's have a view that displays the country name rather than code.

```
create view vmovies
as select m.movieid,
        m.title,
        m.year_released,
        c.country_name
from movies m
inner join countries c
on c.country_code = m.country
```

from movies *from countries*

movieid	title	year_released	country_name
2	Blade Runner	1982	United States
6	Das Boot	1985	Germany
9	Goodfellas	1990	United States
15	Le cinquième élément	1997	France
22	The Lord of the Rings	2001	New Zealand
27	We Feed the World	2005	Australia
25	Ying hung boon sik	1986	Hong Kong



Never understood why people were confusing both.

Wrong!
AUSTRIA,
not Australia!

CORRECTION SQL Server would let you update ... and try to change the name in table COUNTRIES.

```
create view vmovies
as select m.movieid,
       m.title,
       m.year_released,
       c.country_name
from movies m
     inner join countries c
       on c.country_code = m.country
```


NOT this

THIS

Most products will express concern and prevent you from doing it.

Abandon all hope, ye who enter here

In many cases, view update is simply impossible.



Most joins
Aggregates
Expressions
Omitted mandatory columns (insert)

Sometimes it works very well

In some cases, view update is quite possible.

This will work fine with Oracle, which would have complained with a join

One table
`create or replace view vmy_movies
as select m.movieid,
 m.title,
 m.year_released,
 m.country
from movies m
where m.country in
 (select c.country_code
 from countries c
 inner join user_scope u
 on u.continent = c.continent
 where u.username = user)`

USER_SCOPE

Username	Continent
HUIZHONG	ASIA
PAVEL	EUROPE
IBRAHIM	AFRICA
AMINATA	AFRICA
MICHAEL	EUROPE
JUAN_CARLOS	AMERICA
SANDEEP	ASIA
PATRICIA	AMERICA
PATRICIA	EUROPE


Everything else in a subquery

Which proves that in some cases join and subquery aren't exactly equivalent ...

There is no problem because the view update maps to a simple table update.

**Plain insert/
update/delete
of *movies***

Now, there may STILL be a problem.



Suppose that you are in charge of Asia/Pacific, and only see films from this region.

Consistency Issue

`select * from vmy_movies;`

movieid	title	year_released	country
19	Pathar Panchali	1955	in
20	Shichinin no Samurai	1954	jp
21	Sholay	1975	in
22	The Lord of the Rings	2001	nz
25	Ying hung boon sik	1986	hk
26	We Feed the World	2005	au

Oops

*Only from
Asia/Oceania*

If you change the country from Australia to Austria (in Europe), poof! you no longer see it.

```
update vmy_movies
set country = 'at'
where movieid = 26
```



Nothing prevents from

```
insert into vmy_movies(title, year_released, country)
values ('Snow White and the Seven Dwarfs', 1937, 'us')
```

UNLESS

There is one special constraint, though, that exists for views: **WITH CHECK OPTION**.

```
create or replace view vmy_movies
```

```
as select m.movieid,
```

```
    m.title,
```

```
    m.year_released,
```

```
    m.country
```

```
from movies m
```

```
where m.country in
```

```
    (select c.country_code
```

```
    from countries c
```

```
    inner join user_scope u
```

```
        on u.continent = c.continent
```

```
    where u.username = user)
```

```
with check option
```

It prevents you from making a change that will make a row disappear from the view (other than a DELETE)

CHECK OPTION would let you update from Australia to any Asian country



But not to a country from another region



Solution in some cases:

insert procedure




update procedure



delete procedure


If updating the view directly is impossible, in many cases (remember when we were displaying the country name) what should be applied to base tables is fairly obvious and can be performed by dedicated stored procedures.

**View**

There is a special type of trigger called an

instead of trigger

It can be created on a view and lets you call a procedure "instead of" performing the triggering event




insert procedure

update procedure

delete procedure

This can be quite useful with web development frameworks. Map your framework objects to views with "instead of" triggers, and let the framework generate the maintenance screens for you.

Development frameworks

But bad for massive loads

Sometimes another option than views

Function returning a table

They aren't supported by all DBMS products.

Views are cleaner ...

The advantage of a function returning a table is that it allows you to inject a parameter deep inside a query: with a view, it's just like with a table, you can only apply conditions to the data it returns. No such restriction with a function returning a table, but they are more in the "useful dirty trick" category, and once again every DBMS isn't supporting them, or supporting them in the same way (some products may return "streamed" data, returning rows as they are retrieved, some products may buffer them, retrieving data in memory before returning it). Dangerous if you ever migrate to another DBMS.

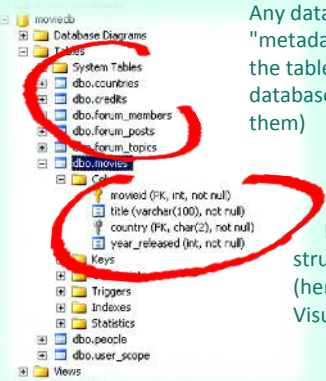
One very good example of view application is the set of tables that contain information about the objects in the database, collectively known as the

Data Dictionary or sometime called the Catalog

They are using all the features we have seen (you only have privileges to read views and only see what is relevant to your account)

One catalog per database

You always have ONE catalog per database. A database is an independent unit and you can have foreign keys only within (inside) one database; however, you can have several schemas in a database, and you can reference tables in another schema. There may also be metadata such as user accounts that is shared among databases. Most DBMS products can manage several databases at once; other than SQLite, the exception is MySQL that only has ONE catalog. What MySQL calls a *database* is actually a schema.



Any database stores "metadata" that describes the tables in your database (and not only them)

All client tools use this information to let you browse the structure of your tables (here it's SQL Server, Visual Studio)

CREATE
DROP
ALTER
GRANT
REVOKE

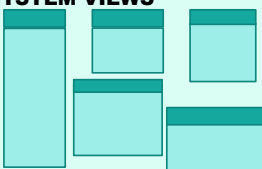
~~INSERT~~
~~DELETE~~
~~UPDATE~~

SYSTEM TABLES

DANGER
KEEP OUT

Whenever you are issuing DDL commands, you are actually modifying system tables. They must NEVER be directly changed.

SYSTEM VIEWS



SQLite **sqlite_master**

Read access to these tables is provided through system views.

information_schema
sysibm **IBM DB2**

+ views in schema **sys** **SQL Server**
+ views in schema **syscat** **IBM DB2**
+ **pg...** views **PostgreSQL**

INFORMATION_SCHEMA.TABLES

In these views you only see what YOU are allowed to see. Only administrators see everything.

SYSCAT.TABLES **IBM DB2**

USER_TABLES / ALL_TABLES **ORACLE**

Those you have created Those you can query

*standard***INFORMATION_SCHEMA****= minimum ...**

The "SQL standard" defines a schema for the catalog that several DBMS vendors try to implement (Oracle, so far, doesn't follow it, perhaps because Oracle has no schemas independent from user accounts, and DB2 doesn't call it INFORMATION_SCHEMA). However, you only find minimum information in INFORMATION_SCHEMA. Some products have views to describe triggers, others haven't, for instance. Other than a small common set, many columns may also be different simply because implementations are different.

```

moviedb-> select table_name,
moviedb->           column_name,
moviedb->           ordinal_position,
moviedb->           data_type
moviedb-> from information_schema.columns
moviedb-> where table_name = 'movies';

```

table_name	column_name	ordinal_position	data_type
movies	year_released	4	integer
movies	country	3	character
movies	title	2	character varying
movies	movieid	1	integer

There are usually simpler commands to display the structure of a table, but these commands execute nothing more than this type of query. Everything is pulled out of the data dictionary. This MySQL query gives the same result on PostgreSQL (very rare!).

INFORMATION

As a developer, you can get from the data dictionary some information that is hard to get elsewhere (constraints, for instance). Database administrators use them a lot for scripting, because the data dictionary always reflects the current state of a database.

SCRIPTING

(DBA )

```

select 'drop table ' ||
       table_name || ';'
from information_schema.tables
where table_name like 'TMP%'
and ...

```

DBAs often use queries on the catalog to generate other SQL queries; it can be done in a script, or in a procedure with a cursor (a case when cursors are mandatory). They sometimes generate other commands, such as shell script, for instance for backing up database files (the name of which can be found in some remote corners of the catalog).