

CS307

Database Principles

Stéphane Faroult
sfaroult@roughsea.com

朱悦铭 Zhu Yueming zhuym@sustc.edu.cn

We have seen that if the idea of "data warehousing" was first introduced by Bill Inmon, many original ideas were contributed by Ralph Kimball.



Ralph Kimball



Bill Inmon

In particular the idea of "facts" and heavily denormalized "dimensions".

We'll store date hierarchies that will ease aggregates

Dimension table

Hierarchy

FullDate
DayOfWeek
DayNumInMonth
DayNumOverall
DayName
DayAbbrev
WeekDayFlag
WeekNumInYear
WeekNumOverall
WeekBeginDate
WeekBeginDateKey
Month
MonthNumOverall
MonthName
Quarter
Year
YearMonth
FiscalMonth
FiscalQuarter
FiscalYear
LastDayInMonthFlag
SameDayYearAgo


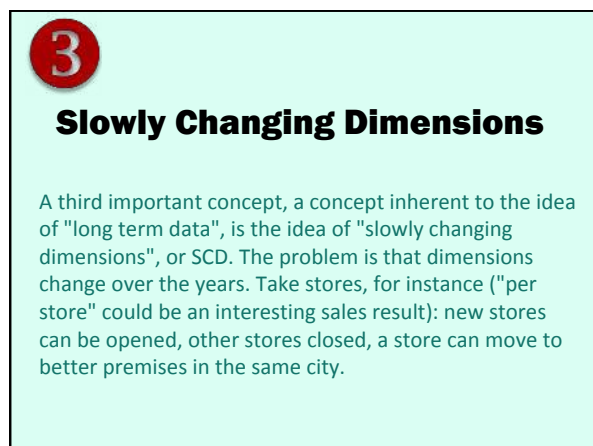
Year
↓
Quarter
↓
Month
↓
Week
↓
Day

FullDate
DayOfWeek
DayNumInMonth
DayNumOverall
DayName
DayNameInSpanish
DayNameInFrench
DayNameInRussian
DayNameInArabic
DayAbbrev
DayAbbrevInSpanish
DayAbbrevInFrench
DayAbbrevInRussian
DayAbbrevInArabic
WeekDayFlag
WeekNumInYear
WeekNumOverall
WeekBeginDate
WeekBeginDateKey
Month
MonthNumOverall
MonthName

Dimension table

We can even adapt to multinational environments and do what we should never do in a decent database.

But then ANYBODY will be able to query with decent performance.




B2B


If you are in a "business to business" (B2B) type of company, some of your customers/suppliers may merge or be taken over by others. You may need to still identify them as the same companies yet.

M&A


Where is my customer/supplier?

YEAR N

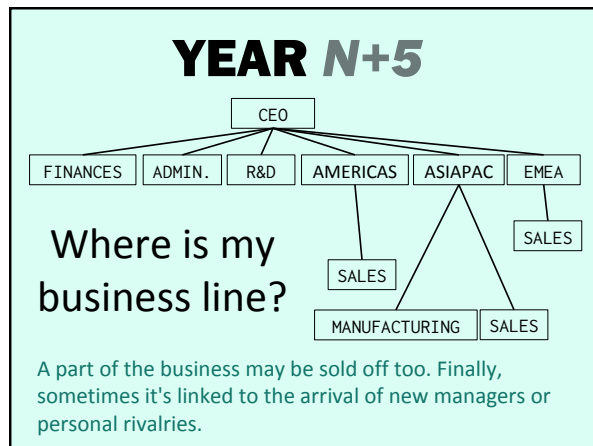
Most very big companies are in a perpetual state of reorganization.

YEAR N+1

It may be linked to the recognition of the importance of one particular department, or the development of new markets.

YEAR N+3

It may be linked to changing business conditions, to new management fashions.



Kimball has tried to classify how to handle slowly changing dimensions. One way to handle them is to ignore them (iPhone 25 same product as iPhone 1)

Type 0 SCD

Keep as it was first inserted

Or you can update, and consider that there is Russia and there never was such thing as Soviet Union. Or consider that there is BDR and DDR, and there is Germany. Or that BDR and DDR became Germany in 1990.

Type 1 SCD

Update dimension - lose old record

Type 2 SCD

Add a new record — Disconnect old facts/new facts

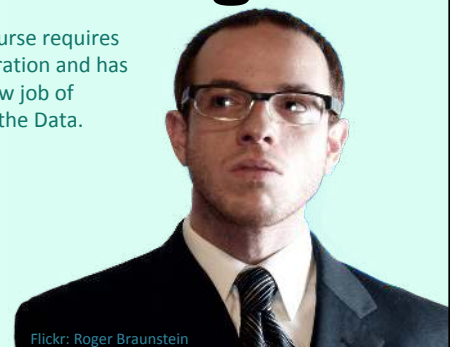
Type 3 SCD

In practice most SCDs are type 1 or 2

Add old/new/date information to the dimension

Data Manager

All this of course requires due consideration and has created a new job of Guardian of the Data.



Flickr: Roger Braunstein

THE PLUMBING

Now, we have seen that in all likelihood most queries on a data warehouse will be massive aggregates, that take time. Worst, they will be run by people high up in the hierarchy, whose time is precious (especially when you consider their salaries) and who, having signed the checks and knowing how much all this did cost, expect bang for their bucks.

How can we have blazingly fast aggregates? With a traditional DBMS (there are other options, remember the columnar database presented by Guy Harrison) there are basically two solutions

Speeding up aggregates

Pre-computing aggregates

NOT mutually exclusive ...

When we want a query that says, for instance, "what was the sum of sales per province and month last year", from an SQL standpoint getting good performance is tough. We'll return at most 22x12 rows, which is peanuts (it won't be much better with municipalities, autonomous regions and special administrative regions). We'll scan a small table of geographical areas, and a small table of dates. Neither "province" nor "month" will be very selective.


Because of 2NF, cardinality is often low

The only restriction "last year", is rather weak

Only broad criteria

And we may aggregate MANY rows.

Dimensions to facts



Very often, you go from referencing table to referenced table when you join. Here, all criteria will systematically be applied to reference tables that DON'T point to the facts table (it's the opposite)

For low cardinality columns, advanced products such as Oracle implements "bitmap indexes". Their entries are chunks of bits that cover many rows and contain a bit that says whether the row contains the key or not.

Bitmap indexes

MON	1000000100000010000001000000100000
TUE	0100000010000001000000100000010000
WED	0010000001000000100000010000001000
THU	0001000000100000010000001000000100
FRI	0000100000010000001000000100000010
SAT	0000010000001000000100000010000001
SUN	0000001000000100000010000001000000

The advantage of bitmap indexes is that by applying AND and OR operations to the bitmaps you can quickly identify the rows that satisfy several conditions (Shenzhen and November) and fetch them.

Bitmap indexes

11110010110000010000111010001000111
 0100010100100010001110010011110000
 0011100010011001000100101100010010
 1001011000101100000010001110000100

AND
OR

Facts Dimensions

Star transformation ORACLE

```
SELECT c.cust_city, t.calendar_quarter_desc,
       SUM(s.amount_sold) sales_amount
FROM sales s, times t, customers c, channels ch
WHERE s.time_id = t.time_id
      AND s.cust_id = c.cust_id
      AND s.channel_id = ch.channel_id
      AND c.cust_state_province = 'CA'
      AND ch.channel_desc = 'Internet'
      AND t.calendar_quarter_desc IN ('1999-01', '1999-02')
GROUP BY c.cust_city, t.calendar_quarter_desc
```

Oracle also implements a parameter that enables the "star schema transformation"

This transformation is just adding redundant conditions that make the choice of the suitable plan irresistible to the optimizer.

Star transformation ORACLE

```
SELECT c.cust_city, t.calendar_quarter_desc,
       SUM(s.amount_sold) sales_amount
FROM sales s, times t, customers c, channels ch
WHERE s.time_id = t.time_id
      AND s.cust_id = c.cust_id
      AND s.channel_id = ch.channel_id
      AND c.cust_state_province = 'CA'
      AND s.cust_id IN (SELECT cust_id
                        FROM customers
                        WHERE cust_state_province='CA')
      AND ch.channel_desc = 'Internet'
      AND t.calendar_quarter_desc IN ('1999-01', '1999-02')
GROUP BY c.cust_city, t.calendar_quarter_desc
```

Star transformation ORACLE

These redundant conditions are added for every dimension in the query

Similar idea, hand-crafted

We only need to hit the index on sales(cust_id)

```
SELECT s.rowid
FROM sales s, customers c
WHERE s.cust_id = c.cust_id (it contains cust_id and rowid)
AND c.cust_state_province = 'CA'

INTERSECT
SELECT s.rowid
FROM sales s, times t
WHERE s.time_id = t.time_id
AND t.calendar_quarter_desc IN ('1999-01', '1999-02')

INTERSECT
SELECT s.rowid
FROM sales s, channels ch
WHERE s.channel_id = ch.channel_id
AND ch.channel_desc = 'Internet'
```

"rowid" is Oracle-speak for "row address". You can try to identify the rows you need first.

We are assuming one separate index per FK

Similar idea, hand-crafted

Join to sales, and aggregate

Once you know exactly which rows you want, you can prey on them in the big table.

Database Machines/Appliances

TERADATA.

NETEZZA
an IBM Company

Greenplum

Oracle
ExadataX3

and others ...

Several companies have created dedicated hardware that is more specially geared towards the data warehouse type of queries. Some databases, called columnar databases, store data by columns rather than rows to speed up aggregates.



Precomputing aggregates

Another, simple way to speed up aggregates is NOT to compute them on the fly but, if we have an idea about the aggregates that will be computed most often, to precompute them every time we upload the database, at least partially.

"Materialized views", which are actually tables, are saved query results (very often aggregates) that are prescheduled and automatically refreshed. You can query the materialized view directly, but some optimizers are smart enough to recognize when they can use a pre-computed result in a query on the original table. It's OK on data that isn't "real time data".

Materialized Views

NOT views

View = saved query text

Materialized view = saved query result

You specify in the creation statement the refresh schedule. It can be a full cancel and replace or, sometimes, an incremental refresh that requires a trigger-populated log on the original table(s).

Materialized Views

Refresh schedule

Full refresh `insert into ...
select ...`


Fast refresh



Of course, it's all plain replicated data.

Some products go further and try to compute every possible aggregate possible, detail AND all possible sums.

Cubes



If you don't ask for anything that strays from the predefined path, you get immediate results.

Cubes





Image found on slidewiki.org

Cubes

Things become funny when you have MANY dimensions – the number of combinations explodes ...



FUNDAMENTALS

Database Design

Indexing

There are two basic pre-conditions for good performance: a proper database design, and proper indexing. A bad design leads to very hard queries, hard to write, and hard for the optimizer to improve. Proper constraints provide you with the main indexes required by joins.

Test: 5 million rows

Rows returned	Speed improvement with index
~ 1,000,000	x1.1
~ 500,000	x1.0
~ 50,000	x1.0
~5,000	x1.1
~1,000	x20.8

Keep in mind that the real benefit of index is only when you return very few rows. The measurements above were taken with MySQL, with Oracle and SQL Server scans perform even better.

If you write queries in a way that prevents from using indexes ...

... options will be limited for the optimizer


The optimizer can recombine a query, it can switch between join, correlated/uncorrelated subquery, it cannot rewrite functions called at the wrong place or automatic data type conversions that return a result, but prevent from descending indexes. Badly writing queries is like closing doors.

The main problem of the optimizer is getting an estimate of how many rows are returned at each step when processing a query.


Selectivity?

How many rows returned at each step?

(cardinality estimate)



Flickr: Brian Hilligas



It won't be able to compute it and decide about using or not using indexes (few rows, indexes good, many rows, indexes bad) if it hasn't statistics. But what kind of statistics?

People



1,300,000,000



60,000,000

In terms of population size, there is no contest between China and France.

Family Names



3,000




1,300,000


BUT when it comes to surnames, it's the opposite. France and Italy are the countries with the most different surnames. It's common for French schoolboys (more than girls) to call each other by their surnames only, it sometimes happens in companies too.

One Family Name






400,000 



46 




It means that on average, there are almost half a million Chinese per surname, and under 50 French. But averages are misleading.

The number of Lis in China is about the same as the population of the Philippines (12th most populated country in the world)

李 > 100,000,000 
 王 > 95,000,000 
 张(張) > 90,000,000 

The numbers of Wangs and Zhangs are pretty impressive too.

Even in France, some surnames are extremely common. And the top one is common in the UK and in the US too (although not pronounced the same)

Martin > 230,000 
 Bernard > 100,000 
 Dubois > 90,000 

Other surnames look different only because of spelling variations (Lefebvre, Lefèvre, Lefébure are basically the same name – and mean the same as "Smith").


Basic statistics on columns are composed of minimum and maximum values, as well as the number of distinct values; but as we have just seen with surnames, dividing the number of rows in a table by the number of distinct values in a column (assuming uniform distribution) gives a pretty wrong idea of how many rows will be associated with one value.

Minimum value

Maximum value

Number of distinct values

We need a better idea of a distribution that may be bell-shaped, exponential, or absolutely anything with discrete values.



What does the optimizer know about data?

We are going to see what a product such as Oracle knows about data values in a column (which can be indexed or not; usually detailed statistics are collected over indexed columns)

```
SQL> create index movie_country_index on movies(country);
```

Index created.

```
SQL> begin
2   dbms_stats.gather_table_stats(ownname=>user,
3                                   tabname=>'movies',
4                                   cascade=>true,
5                                   method_opt=>'for all columns size auto');
6 end;
7 /
```

PL/SQL procedure successfully completed.

I'm creating an index on countries on table MOVIES in a subset of the film database, and calling the Oracle procedure that collects statistics on a table.

```
SQL> select table_name, num_rows, blocks, avg_row_len
2   from user_tables
3  where table_name = 'MOVIES';
```

TABLE_NAME	NUM_ROWS	BLOCKS	AVG_ROW_LEN
MOVIES	183	5	26

SQL>

Basic stats about the table are the number of rows, how many data blocks are used to physically store the table (it doesn't take indexes into account) and the size of the average row.

```
SQL> select table_name, num_rows, blocks, avg_row_len
2   from user_tables
3  where table_name = 'MOVIES';
```

TABLE_NAME	NUM_ROWS	BLOCKS	AVG_ROW_LEN
MOVIES	183	5	26

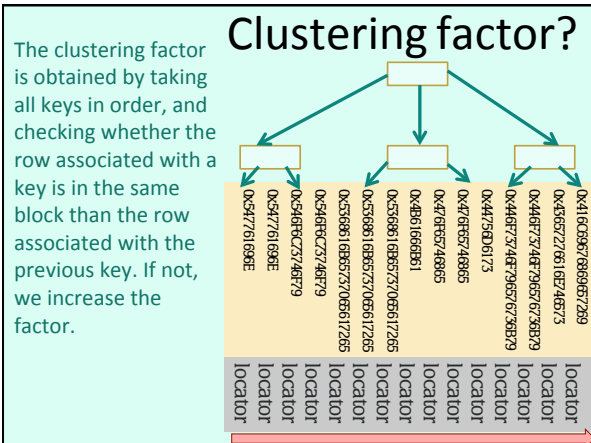
```
SQL> select index_name, uniqueness, distinct_keys, clustering_factor
2   from user_indexes
3  where table_name = 'MOVIES';
```

INDEX_NAME	UNIQUENES	DISTINCT_KEYS	CLUSTERING_FACTOR
SYS_C0012408	UNIQUE	183	1
SYS_C0012409	UNIQUE	183	1
MOVIE_COUNTRY_INDEX	NONUNIQUE	37	1

For indexes, uniqueness, keys and "clustering factor" (indexes named SYS_Cxxx are associated with unnamed primary key and unique constraints)

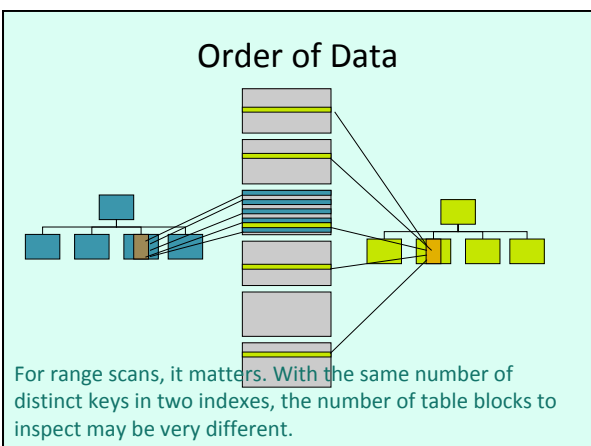
Distinct keys:
one composite index
better than
several single-column indexes

But remember that if there is no condition on the first column in the index ...



Relationship between order of rows and order of keys

What is the result? If the clustering factor is close to the number of rows in the table, that means that two successive keys are generally in different blocks. If it's close to the number of blocks in the table, it means that the rows are more or less in the same order as the keys.



Lower clustering factor



Fewer blocks
to fetch and inspect

Columns?

We have seen what the Oracle optimizer knows about tables, and what it knows about indexes. What can it know about column values?

```
SQL> select column_name, data_type,
2      nullable, num_distinct,
3      low_value, high_value, density,
4      num_nulls, num_buckets
5 from user_tab_columns
6 where table_name = 'MOVIES'
7 order by column_id;
```

Column name, data type,
nullable and num_nulls
(number of nulls) require
no comment.

COLUMN_NAME	DATA_TYPE	N	NUM_DISTINCT	LOW_VALUE
HIGH_VALUE	DENSITY	NUM_NULLS	NUM_BUCKETS	
MOVIEID	NUMBER	N	183	C102
C20254	.005464481	0	1	
TITLE	VARCHAR2	N	174	3132207374756C7
5A68656E2078696	.00273224	0	174	
COUNTRY	CHAR	N	37	6172
7573	.00273224	0	37	
YEAR_RELEASED	NUMBER	N	69	C21416
C2150A	.00273224	0	69	

```
SQL> select column_name, data_type,
2      nullable, num_distinct,
3      low_value, high_value, density,
4      num_nulls, num_buckets
5 from user_tab_columns
6 where table_name = 'MOVIES'
7 order by column_id;
```

The number of distinct
values gives a good idea of
selectivity, even when not
unique (TITLE)

COLUMN_NAME	DATA_TYPE	N	NUM_DISTINCT	LOW_VALUE
HIGH_VALUE	DENSITY	NUM_NULLS	NUM_BUCKETS	
MOVIEID	NUMBER	N	183	C102
C20254	.005464481	0	1	
TITLE	VARCHAR2	N	174	3132207374756C7
5A68656E2078696	.00273224	0	174	
COUNTRY	CHAR	N	37	6172
7573	.00273224	0	37	
YEAR_RELEASED	NUMBER	N	69	C21416
C2150A	.00273224	0	69	

```
SQL> select column_name, data_type,
2      nullable, num_distinct,
3      low_value, high_value, density,
4      num_nulls, num_buckets
5 from user_tab_columns
6 where table_name = 'MOVIES'
7 order by column_id;
```

Minimum and maximum values are encoded, "density" assumes uniform distribution.

COLUMN_NAME	DATA_TYPE	N	NUM_DISTINCT	LOW_VALUE
HIGH_VALUE	DENSITY	NUM_NULLS	NUM_BUCKETS	
MOVIEID	NUMBER	N	183	C102
C20254	.005464481	0	1	
TITLE	VARCHAR2	N	174	3132207374756C7
5A68656E2078696	.00273224	0	174	
COUNTRY	CHAR	N	37	6172
7573	.00273224	0	37	
YEAR_RELEASED	NUMBER	N	69	C21416
C2150A	.00273224	0	69	

```
SQL> select column_name, data_type,
2      nullable, num_distinct,
3      low_value, high_value, density,
4      num_nulls, num_buckets
5 from user_tab_columns
6 where table_name = 'MOVIES'
7 order by column_id;
```

More interesting, num_buckets is what Oracle uses for tracking distribution with histograms.

COLUMN_NAME	DATA_TYPE	N	NUM_DISTINCT	LOW_VALUE
HIGH_VALUE	DENSITY	NUM_NULLS	NUM_BUCKETS	
MOVIEID	NUMBER	N	183	C102
C20254	.005464481	0	1	
TITLE	VARCHAR2	N	174	3132207374756C7
5A68656E2078696	.00273224	0	174	
COUNTRY	CHAR	N	37	6172
7573	.00273224	0	37	
YEAR_RELEASED	NUMBER	N	69	C21416
C2150A	.00273224	0	69	

Histograms:

Useless if unique

If a column is unique, the number of buckets will be one or zero, meaning no histogram.

More of fewer than 256 distinct values?

If there are fewer than 256 values, Oracle will count how often each one occurs, and each counter will be a "bucket".

Frequency Histogram

When each individual value is counted, this is called a "frequency histogram" and the result is of course extremely precise.

This is what Oracle holds in a dictionary table.

```
SQL> select column_name, endpoint_number, endpoint_value, endpoint_actual_value
2   from user_tab_histograms
3   where table_name = 'MOVIES'
4   order by 1, 2;
```

COLUMN_NAME	ENDPOINT_NUMBER	ENDPOINT_VALUE	ENDPOINT_ACTUAL_VALUE
COUNTRY	2	5.0597E+35	
COUNTRY	3	5.0601E+35	
COUNTRY	5	5.1090E+35	
COUNTRY	6	5.1094E+35	
COUNTRY	11	5.1116E+35	
COUNTRY	13	5.1601E+35	
COUNTRY	15	5.1615E+35	
...			
COUNTRY	112	5.8371E+35	
COUNTRY	114	5.8373E+35	
COUNTRY	115	5.9418E+35	
COUNTRY	127	5.9430E+35	
COUNTRY	130	5.9917E+35	
COUNTRY	131	5.9921E+35	
COUNTRY	134	5.9939E+35	
COUNTRY	135	6.0472E+35	
COUNTRY	183	6.0983E+35	

RAW

This is what it means ("end points" are cumulative).

Number of rows in the table: 183 (last computed ...)
Sample size 100.000% (183 rows) collected ...

ENDPOINT	POPULARITY	FREQUENCY
ar		2 =
at		1
be		2 =
bg		1
br		5 =
...		
fr		7 =
gb		6 =
hk		7 =
id		1
in		26 =
it		9 =
jp		7 =
...		
pl		2 =
ro		1
ru		12 =
se		3 =
sg		1
sp		3 =
tw		1
us		48 =

COOKED

Height-Balanced Histogram (Oracle <= 11gR2)

When there are more than 255 different values, Oracle will use 254 buckets but will use a different algorithm to distribute values to the bucket. Here is what it was doing up to version 12.

Assume that the limit is 5
instead of 254 "buckets"

We'll do it for countries but assume that we want to distribute them between a far smaller number of buckets than what Oracle would do.

```
SQL> select country, count(*) from movies
2 group by country
3 order by country;
```

CO	COUNT(*)
ar	2
at	1
be	2
bg	1
br	5
ca	2
ch	2
cn	6
cz	2
de	5
dk	3
eg	5
fr	7
gb	6
hk	7

Here is the number of countries,
with the number of films for each.

183 rows - divide by 5 buckets

=> 37 rows by bucket

If we want to distribute them across 5 buckets, then
we want 37 rows per bucket.

```
SQL> select country, count(*) from movies
2 group by country
3 order by country;
```

CO	COUNT(*)
ar	2
at	1
be	2
bg	1
br	5
ca	2
ch	2
cn	6
cz	2
de	5
dk	3
eg	5
fr	7
gb	6
hk	7

We are taking the list from the
beginning, and as long as we
haven't 37 rows we assume they all
go to the first bucket. The first
bucket fills up with the first French
film, so 'fr' indicates the end of the
first bucket.

=> 1 in first bucket, 6 in next

hk	7
id	1
in	26 => 17 in 2 nd bucket, 9 in next
it	9
jp	7
kr	1
ma	2
mx	2
my	1
ne	1
ng	1
nl	2
no	1
nz	1
pk	1
pl	2
ro	1
ru	12
se	3
sg	1
sp	3
tw	1
us	48

And we go on, filling buckets one
after the other, distributing equally
(give or take one) between
buckets.

hk	7	
id	1	
in	26	=> 17 in 2 nd bucket, 9 in next
it	9	
jp	7	
kr	1	
ma	2	
mx	2	
my	1	
ne	1	
ng	1	
nl	2	
no	1	
nz	1	
pk	1	=> end of 3 rd bucket
pl	2	
ro	1	
ru	12	
se	3	
sg	1	
sp	3	
tw	1	
us	48	=> end of 4 th and 5 th buckets
SQL>		

Bucket	end value
1	fr
2	in
3	pk
4	us
5	us

"bucket boundaries" are supposed to be "popular" (i.e. frequent) values. Obviously wrong for pk and we missed ru.

Hybrid Histogram

(Oracle >= 12c)

Don't split values between buckets

The height-balanced histograms were often giving less than convincing results, and many DBAs were preferring having no histograms rather than height-balanced histograms. This is why the algorithm has been changed in Oracle 12.

SQL> select country, count(*) from movies		
2 group by country		
3 order by country;		
CO	COUNT(*)	
ar	2	
at	1	
be	2	
bg	1	
br	5	
ca	2	
ch	2	
cn	6	
cz	2	
de	5	
dk	3	
eg	5	
fr	7	=> 43 values in first bucket
gb	6	
hk	7	

With hybrid histograms, even if the goal is still to have roughly 37 countries per bucket, we'll cram all the films from one country in one bucket and won't split them between two buckets.

hk	7	
id	1	
in	26	=> 40 values in 2 nd bucket
it	9	
jp	7	
kr	1	
ma	2	
mx	2	
my	1	
ne	1	
ng	1	
nl	2	
no	1	
nz	1	
pk	1	
pl	2	
ro	1	
ru	12	=> 44 values in 3 rd bucket
se	3	
sg	1	
sp	3	
tw	1	
us	48	=> 56 values in 4 th bucket
SQL>		

Bucket	end value	size
1	fr	43
2	in	40
3	ru	44
4	us	48

The preceding anomaly is gone.

Better but not perfect ...

This said, you will notice that even if there are MORE Italian and AS MANY Hong-Kong films as French ones, neither "it" nor "hk" appears as a popular value. We may still have cases when the optimizer fails to estimate correctly how many rows will be associated with a country value. Of course, in a complex query poor estimates may cascade.

Histogram issues

Far away values

Oracle doesn't index nulls

Another problem with Oracle is linked to something which is common with dates, the use of dummy values. As nulls may not be stored in Oracle and may have no address, they cannot be indexed. Columns "valid_until" are frequent, and rather than having a null to indicate current values, people often prefer using a dummy, far away date that can be indexed.

Histogram issues

Far away values

~~31st Dec 9999~~

Except that using the maximum date allowed by Oracle is a bad idea. If your oldest date is 2000, asking for everything until 01/01/2018 will probably mean returning most of the rows. The optimizer may see 18 years out of a range of 8000 (remember, it knows min and max) and think that it's a small percentage that would be fetched faster with an index.

Histogram issues

False range scan

Another issue which isn't a histogram issue proper but for which you might think that histograms would help when they probably won't is what I call a "false range scan".

Histogram issues

False range scan

Real where event_date between sysdate -7 and sysdate

↑ ↑ ↑
column constant constant

A true range scan is when you ask for values in a column that are between two constant (or computed) values. With histograms, the optimizer can get a fairly good idea of the number of rows returned, and great if the index has a good clustering factor.

Histogram issues

False range scan

where event_date between sysdate -7 and sysdate

False where sysdate between start_date and end_date

↑ ↑ ↑
constant column1 column2


A false range scan looks deceptively like a real one, except that here you are comparing ONE constant to TWO columns instead of the reverse.

Histogram issues

False range scan

where event_date between sysdate -7 and sysdate

where sysdate >= start_date
and sysdate <= end_date



What you are really doing is taking the intersect of two open-ended intervals. And here, getting an estimate of the number of rows is FAR more difficult, and indexes won't help much.

Histogram issues

False range scan

ip_range_start ip_range_end location

-----+-----+-----

valid_credit_card_start valid_credit_card_end

-----+-----

It's a problem that is commonly encountered, as soon as you have a "from" column and a "to" column.

Finally, the optimizer may far overestimate a number of rows returned because of correlations between columns.

Correlation

The optimizer assumes 3rd Normal Form

... attributes only depend on the key and are independent

However, saying "independent" is often a bit far fetched.

Not completely independent

	date_start	date_end	

If I tell you "how many people where born after 1980 and died before 1940" you'll tell me immediately 0. For the optimizer, if 20% of people were born after 1980 and 20% died before 1940, it will be $20\% \times 20\% = 4\%$ of the total number of rows. You can ask Oracle to compute stats on groups of columns.

Design and Statistics

"Artificial" Correlations

Knowing which statistics to collect and, from there, trying to estimate how many rows should be returned by such or such parameter value (a crucial step in the choice of an execution plan) is hard enough for the DBMS system. Poor design choices make it far more difficult for the optimizer and increase significantly the chances of its taking a bad decision.

Let's take for example the "Entity/attribute/Value" model so loved of people who start coding without really knowing which attributes their data will ultimately hold.

PEOPLE

entity	attribute	value
567	first_name	Gene
567	surname	Tierney
567	birth	1920
567	death	1991
567	gender	F

What should be horizontal is vertical. Not only can you forget about foreign keys, but the optimizer will struggle.

468 rows

346 different first names (456 people have one)
418 different surnames
106 different birth years
71 different death years (for 237 people)
2 genders

I have computed a few statistics of my own on a relatively small sample of actors and directors. First I stored my data in a traditional way (id, first name, surname, born, died, gender)

Entity/Attribute/Values

5 different attributes – almost uniform (null won't appear)

Rows = surname + born + gender
 for everybody **1404**
 first_name for almost everybody **456**
 death year for half the people **237**
Total **2097** rows

With the EAV model, the same data becomes a table with half the number of columns and more than three times the number of rows. Only known values are stored.

With a standard model, how many rows will be returned by a condition on gender is easy to compute.

gender = 'F' ~50%, ~230 rows

With an EAV model, because the attribute name is now part of the data, assuming data independence is wrong. 'Gender' is always associated with either 'M' or 'F'.

attr_name = 'Gender' 1/5 20%
and attr_value = 'F' 230/2097 11%

20% * 11% ?

Assuming that columns are independent, the optimizer will believe that only 2% of the 2,000 rows will be returned, and will be wrong by an order of magnitude.

High odds to pick the wrong algorithm

As, if you remember Kyte's allegory, the suitability of an algorithm depends on data volumes, an EAV model will underestimate volumes and pick on indexes and algorithms that use indexes.

"Extended statistics" on several columns

Hence the benefit of "extended statistics" (which only the top products such as Oracle allow), not only to detect genuine correlation in the data (such as between "born" and "died") but also to try to correct artificial correlation introduced by the data model (because any value cannot be associated with any attribute).



Bind variables

Taken into account during parsing

Stability issues come from bind variables that have a distribution value that isn't uniform. You'll never have a problem with ids or unique columns, just with columns in which some values are very common, and others fairly rare.

```
select ...  
from ...  
where datecol between  
      cast('2013-12-01' as date)  
and cast('2013-12-31' as date)  
and status = 'COMPLETED'
```



The problem may be more obvious when the column is indexed, but in fact it's more general than that because even a non-indexed column can be important for estimating the size of the result-set.


```
select ...
from ...
where datecol between
    cast('2013-12-01' as date)
and cast('2013-12-31' as date)
and status = 'ERROR'
```



RARE

You may be using the index on datecol to access data, but if this query is a subquery fed into an operation such as a join, it matters to know whether it returns most of the rows between the two dates or very few.

explain

And the index that may be important to use may be a completely different index.



Index used?

```
select ...
from ...
where datecol between
    cast(:date_min as date)
and cast(:date_max as date)
and status = :status
```

Check parameters
when parsing?



As we have already seen, the parameters may be checked when parsing but because parsed queries may age out of a buffer and being reparsed with different parameters, plans may completely change later in the day. The plan is usually adequate for the first execution; it may be completely wrong for some subsequent executions.

UNPREDICTABLE is worse than

SLOW

Unpredictability is something that everybody hates. You can live with something that is consistently mediocre. Something that can be sometimes blazingly fast and sometimes horribly slow for no apparent reason is something you want to avoid.